

CPE 311 - Computational Thinking with Python

Name: Gwyneth D. Esperat

Section: CPE22S3

Date: March 27, 2024

Github Link: [Module 8](#)

8.1.1 Intended Learning Outcomes

After this activity, the student should be able to:

- Demonstrate querying and merging of dataframes
- Perform advanced calculations on dataframes
- Aggregate dataframes with pandas and numpy
- Work with time series data

8.1.2 Resources

- Computing Environment using Python 3.x
- Attached Datasets (under Instructional Materials)

8.1.3 Procedures

The procedures can be found in the canvas module. Check the following under topics:

- 8.1 Weather Data Collection
- 8.2 Querying and Merging
- 8.3 Dataframe Operations
- 8.4 Aggregations
- 8.5 Time Series

```
import pandas as pd

# Load the earthquakes data
earthquakes = pd.read_csv('/content/earthquakes.csv')

# Filter earthquakes in Japan with magType 'mb' and magnitude >= 4.9
earthquakes_japan = earthquakes[(earthquakes['place'].str.contains('Japan')) &
                                  (earthquakes['magType'] == 'mb') &
                                  (earthquakes['mag'] >= 4.9)]

earthquakes_japan
```

	mag	magType	time	place	tsunami	parsed_place	
	1563	4.9	mb	1538977532250	293km ESE of Iwo Jima, Japan	0	Japan
	2576	5.4	mb	1538697528010	37km E of Tomakomai, Japan	0	Japan
	3072	4.9	mb	1538579732490	15km ENE of Hasaki, Japan	0	Japan
	3632	4.9	mb	1538450871260	53km ESE of Hitachi, Japan	0	Japan

```
# Create bins for magnitudes
bins = pd.interval_range(start=0, end=int(earthquakes['mag'].max()) + 1, freq=1)
earthquakes['magnitude_bin'] = pd.cut(earthquakes[earthquakes['magType'] == 'ml']['mag'], bins=bins)

# Count earthquakes in each bin
magnitude_counts = earthquakes.groupby('magnitude_bin').size()

magnitude_counts
```

magnitude_bin	
(0, 1]	2207
(1, 2]	3105
(2, 3]	862
(3, 4]	122
(4, 5]	2
(5, 6]	1
(6, 7]	0
(7, 8]	0

```
dtype: int64

# Load the FAANG data
faang = pd.read_csv('/content/faang.csv', parse_dates=['date'])
faang.set_index('date', inplace=True)

# Group by ticker and resample to monthly frequency with aggregations
faang_monthly = faang.groupby('ticker').resample('M').agg({
    'open': 'mean',
    'high': 'max',
    'low': 'min',
    'close': 'mean',
    'volume': 'sum'
})

faang_monthly
```

		open	high	low	close	volume
ticker	date					
AAPL	2018-01-31	170.714690	176.6782	161.5708	170.699271	659679440
	2018-02-28	164.562753	177.9059	147.9865	164.921884	927894473
	2018-03-31	172.421381	180.7477	162.4660	171.878919	713727447
	2018-04-30	167.332895	176.2526	158.2207	167.286924	666360147
	2018-05-31	182.635582	187.9311	162.7911	183.207418	620976206
	2018-06-30	186.605843	192.0247	178.7056	186.508652	527624365
	2018-07-31	188.065786	193.7650	181.3655	188.179724	393843881
	2018-08-31	210.460287	227.1001	195.0999	211.477743	700318837
	2018-09-30	220.611742	227.8939	213.6351	220.356353	678972040
	2018-10-31	219.489426	231.6645	204.4963	219.137822	789748068
	2018-11-30	190.828681	220.6405	169.5328	190.246652	961321947
	2018-12-31	164.537405	184.1501	145.9639	163.564732	898917007
AMZN	2018-01-31	1301.377143	1472.5800	1170.5100	1309.010952	96371290
	2018-02-28	1447.112632	1528.7000	1265.9300	1442.363158	137784020
	2018-03-31	1542.160476	1617.5400	1365.2000	1540.367619	130400151
	2018-04-30	1475.841905	1638.1000	1352.8800	1468.220476	129945743
	2018-05-31	1590.474545	1635.0000	1546.0200	1594.903636	71615299
	2018-06-30	1699.088571	1763.1000	1635.0900	1698.823810	85941510
	2018-07-31	1786.305714	1880.0500	1678.0600	1784.649048	97629820
	2018-08-31	1891.957826	2025.5700	1776.0200	1897.851304	96575676
	2018-09-30	1969.239474	2050.5000	1865.0000	1966.077895	94445693
	2018-10-31	1799.630870	2033.1900	1476.3600	1782.058261	183228552
	2018-11-30	1622.323810	1784.0000	1420.0000	1625.483810	139290208
	2018-12-31	1572.922105	1778.3400	1307.0000	1559.443158	154812304
FB	2018-01-31	184.364762	190.6600	175.8000	184.962857	495655736
	2018-02-28	180.721579	195.3200	167.1800	180.269474	516621991
	2018-03-31	173.449524	186.1000	149.0200	173.489524	996232472
	2018-04-30	164.163557	177.1000	150.5100	163.810476	751130388
	2018-05-31	181.910509	192.7200	170.2300	182.930000	401144183
	2018-06-30	194.974067	203.5500	186.4300	195.267619	387265765

```
# Create a crosstab
crosstab_max_magnitude = pd.crosstab(index=earthquakes['tsunami'], columns=earthquakes['magType'],
                                     values=earthquakes['mag'], aggfunc='max')
```

crosstab_max_magnitude

magType	mb	mb_lg	md	mh	m1	ms_20	mw	mwb	mwr	mwW
tsunami										
0	5.6	3.5	4.11	1.1	4.2	NaN	3.83	5.8	4.8	6.0
1	6.1	NaN	NaN	NaN	5.1	5.7	4.41	NaN	NaN	7.5
2018-03-31	1096.108095	1177.0500	980.6400	1091.490476	45430049					

```
# Rolling 60-day aggregations by ticker
rolling_60d = faang.groupby('ticker').rolling(window='60D').agg({
    'open': 'mean',
    'high': 'max',
    'low': 'min',
    'close': 'mean',
    'volume': 'sum'
})
```

rolling_60d

		open	high	low	close	volume
ticker	date					
AAPL	2018-01-02	166.927100	169.0264	166.0442	168.987200	25555934.0
	2018-01-03	168.089600	171.2337	166.0442	168.972500	55073833.0
	2018-01-04	168.480367	171.2337	166.0442	169.229200	77508430.0
	2018-01-05	168.896475	172.0381	166.0442	169.840675	101168448.0
	2018-01-08	169.324680	172.2736	166.0442	170.080040	121736214.0
...
NFLX	2018-12-24	283.509250	332.0499	233.6800	281.931750	525657894.0
	2018-12-26	281.844500	332.0499	231.2300	280.777750	520444588.0
	2018-12-27	281.070488	332.0499	231.2300	280.162805	532679805.0
	2018-12-28	279.916341	332.0499	231.2300	279.461341	521968250.0
	2018-12-31	278.430769	332.0499	231.2300	277.451410	476309676.0

1255 rows × 5 columns

```
# Pivot table for FAANG data
faang_pivot = faang.pivot_table(index='ticker', values=['open', 'high', 'low', 'close', 'volume'], aggfunc='mean')
```

faang_pivot

	close	high	low	open	volume
ticker					
AAPL	186.986218	188.906858	185.135729	187.038674	3.402145e+07
AMZN	1641.726175	1662.839801	1619.840398	1644.072669	5.649563e+06
FB	171.510936	173.615298	169.303110	171.454424	2.768798e+07
GOOG	1113.225139	1125.777649	1101.001594	1113.554104	1.742645e+06
NFLX	319.290299	325.224583	313.187273	319.620533	1.147030e+07

```
# Calculate Z-scores for Netflix data
netflix_data = faang[faang['ticker'] == 'NFLX'].select_dtypes(include=['float64', 'int64'])
```

```
z_scores = netflix_data.apply(lambda x: (x - x.mean()) / x.std())
```

z_scores

	open	high	low	close	volume
date					
2018-01-02	-2.500753	-2.516023	-2.410226	-2.416644	-0.088760
2018-01-03	-2.380291	-2.423180	-2.285793	-2.335286	-0.507606
2018-01-04	-2.296272	-2.406077	-2.234616	-2.323429	-0.959287
2018-01-05	-2.275014	-2.345607	-2.202087	-2.234303	-0.782331
2018-01-08	-2.218934	-2.295113	-2.143759	-2.192192	-1.038531
...
2018-12-24	-1.571478	-1.518366	-1.627197	-1.745946	-0.339003
2018-12-26	-1.735063	-1.439978	-1.677339	-1.341402	0.517040
2018-12-27	-1.407286	-1.417785	-1.495805	-1.302664	0.134868
2018-12-28	-1.248762	-1.289018	-1.297285	-1.292137	-0.085164
2018-12-31	-1.203817	-1.122354	-1.088531	-1.055420	0.359444

251 rows × 5 columns

```
# Create an event dataframe
events_df = pd.DataFrame({
    'ticker': 'FB',
    'date': pd.to_datetime(['2018-07-25', '2018-03-19', '2018-03-20']),
    'event': ['Disappointing user growth announced after close.', 'Cambridge Analytica story', 'FTC investigation']
}).set_index(['date', 'ticker'])
```

```
# Merge with FAANG data
faang_reset = faang.reset_index()
events_merged = pd.merge(faang_reset, events_df, on=['date', 'ticker'], how='outer').set_index(['date', 'ticker'])
```

events_merged

		open	high	low	close	volume	event
date ticker							
2018-01-02	FB	177.68	181.58	177.5500	181.42	18151903	NaN
2018-01-03	FB	181.88	184.78	181.3300	184.67	16886563	NaN
2018-01-04	FB	184.90	186.21	184.0996	184.33	13880896	NaN
2018-01-05	FB	185.59	186.90	184.9300	186.85	13574535	NaN
2018-01-08	FB	187.20	188.90	186.3300	188.28	17994726	NaN
...
2018-12-24	GOOG	973.90	1003.54	970.1100	976.22	1590328	NaN
2018-12-26	GOOG	989.01	1040.00	983.0000	1039.46	2373270	NaN
2018-12-27	GOOG	1017.15	1043.89	997.0000	1043.88	2109777	NaN
2018-12-28	GOOG	1049.62	1055.56	1033.1000	1037.08	1413772	NaN
2018-12-31	GOOG	1050.96	1052.70	1023.5900	1035.61	1493722	NaN

1255 rows × 6 columns

8.1.4 Data Analysis

Provide some comments here about the results of the procedures.

1. Earthquakes in Japan

- **Observation:** A few earthquakes in Japan with a magnitude of 4.9 or greater and magType of mb were identified.
- **Comments:** These specific earthquakes might be significant enough to be studied for patterns or potential impacts on the region, considering their relatively higher magnitude.

2. Magnitude Bins with magType ml

- **Observation:** The majority of earthquakes fall into lower magnitude bins, with a sharp decrease in frequency as magnitude increases.

- **Comments:** This distribution is expected and aligns with the general understanding that lower magnitude earthquakes are far more common than higher magnitude ones.

3. Monthly Aggregations for FAANG Data

- **Observation:** Monthly aggregated data shows the variability in trading volumes and price ranges for each FAANG company over time.
- **Comments:** These trends can help investors understand historical performance, identify seasonal patterns, and make informed decisions.

4. Crosstab between tsunami and magType

- **Observation:** The crosstab revealed the maximum magnitudes associated with tsunamis for different magTypes .
- **Comments:** This analysis could be critical for disaster preparedness and understanding the correlation between earthquake characteristics and tsunami generation.

5. Rolling 60-day Aggregations for FAANG Data

- **Observation:** Rolling averages smooth out short-term fluctuations and highlight longer-term trends in stock prices and trading volumes.
- **Comments:** Investors might use this information to gauge the momentum of a stock and identify potential buying or selling opportunities.

6. Pivot Table for FAANG Data

- **Observation:** The pivot table provided a comparative overview of the average open, high, low, close prices, and volumes for FAANG stocks.
- **Comments:** This comparative analysis is useful for portfolio diversification, showing how different stocks behave over time relative to each other.

7. Z-scores for Netflix's Data

- **Observation:** Calculating Z-scores for Netflix's data highlights how each trading day's prices and volumes deviate from the mean.
- **Comments:** This can help identify outliers or unusual trading days that may be driven by specific events or announcements.

8. FAANG Data with Event Descriptions Added

- **Observation:** Incorporating significant events into the FAANG dataset provides context for stock price movements.
- **Comments:** Understanding the impact of specific events on stock performance is crucial for fundamental analysis and can guide investment strategies.

8.1.5 Supplementary Activity

Using the CSV files provided and what we have learned so far in this module complete the following exercises:

1. With the earthquakes.csv file, select all the earthquakes in Japan with a magType of mb and a magnitude of 4.9 or greater.

```
import pandas as pd

earthquakes = pd.read_csv('/content/earthquakes.csv')
earthquakes
```

	mag	magType	time	place	tsunami	parsed_place
0	1.35	ml	1539475168010	9km NE of Aguanga, CA	0	California
1	1.29	ml	1539475129610	9km NE of Aguanga, CA	0	California
2	3.42	ml	1539475062610	8km NE of Aguanga, CA	0	California
3	0.44	ml	1539474978070	9km NE of Aguanga, CA	0	California
4	2.16	md	1539474716050	10km NW of Avenal, CA	0	California
...
9327	0.62	md	1537230228060	9km ENE of Mammoth Lakes, CA	0	California
9328	1.00	ml	1537230135130	3km W of Julian, CA	0	California
9329	2.40	md	1537229908180	35km NNE of Hatillo, Puerto Rico	0	Puerto Rico
9330	1.10	ml	1537229545350	9km NE of Aguanga, CA	0	California
9331	0.66	ml	1537228864470	9km NE of Aguanga, CA	0	California

9332 rows × 6 columns

```
filtered_earthquakes = earthquakes[
    (earthquakes['place'].str.contains("Japan")) &
    (earthquakes['magType'] == 'mb') &
    (earthquakes['mag'] >= 4.9)
]

print(filtered_earthquakes)
```

	mag	magType	time	place	tsunami	\
1563	4.9	mb	1538977532250	293km ESE of Iwo Jima, Japan	0	
2576	5.4	mb	1538697528010	37km E of Tomakomai, Japan	0	
3072	4.9	mb	1538579732490	15km ENE of Hasaki, Japan	0	
3632	4.9	mb	1538450871260	53km ESE of Hitachi, Japan	0	

parsed_place

1563	Japan
2576	Japan
3072	Japan
3632	Japan

2. Create bins for each full number of magnitude (for example, the first bin is 0-1, the second is 1-2, and so on) with a magType of ml and count how many are in each bin.

```
bins = range(int(earthquakes['mag'].min()), int(earthquakes['mag'].max()) + 2)
mag_ml_df = earthquakes[earthquakes['magType'] == 'ml']

mag_ml_df = mag_ml_df.copy()
mag_ml_df['magnitude_bin'] = pd.cut(mag_ml_df['mag'], bins=bins, include_lowest=True, right=False)

magnitude_counts = mag_ml_df['magnitude_bin'].value_counts().sort_index()

print(magnitude_counts)
```

```
[-1, 0)      446
[0, 1)     2072
[1, 2)     3126
[2, 3)      985
[3, 4)      153
[4, 5)         6
[5, 6)         2
[6, 7)         0
[7, 8)         0
Name: magnitude_bin, dtype: int64
```

3. Using the faang.csv file, group by the ticker and resample to monthly frequency. Make the following aggregations:

- Mean of the opening price
- Maximum of the high price
- Minimum of the low price
- Mean of the closing price
- Sum of the volume traded

```
import pandas as pd

# Corrected: Load the FAANG data
faang_df = pd.read_csv('/content/faang (1).csv', parse_dates=['date'])
faang_df.set_index('date', inplace=True)

# Group by ticker and resample to monthly frequency
faang_monthly = faang_df.groupby('ticker').resample('M').agg({
    'open': 'mean',
    'high': 'max',
    'low': 'min',
    'close': 'mean',
    'volume': 'sum'
})

print(faang_monthly.head())
```

		open	high	low	close	volume
ticker	date					
	2018-01-31	170.714690	176.6782	161.5708	170.699271	659679440
	2018-02-28	164.562753	177.9059	147.9865	164.921884	927894473
	2018-03-31	172.421381	180.7477	162.4660	171.878919	713727447
	2018-04-30	167.332895	176.2526	158.2207	167.286924	666360147
	2018-05-31	182.635582	187.9311	162.7911	183.207418	620976206

4. Build a crosstab with the earthquake data between the tsunami column and the magType column. Rather than showing the frequency count, show the maximum magnitude that was observed for each combination. Put the magType along the columns.

```
earthquakes = pd.read_csv('/content/earthquakes.csv')

# Build a crosstab showing the maximum magnitude for each combination of tsunami and magType
tsunami_magType_max_mag = pd.crosstab(index=earthquakes['tsunami'], columns=earthquakes['magType'],
                                       values=earthquakes['mag'], aggfunc='max')

tsunami_magType_max_mag
```

magType	mb	mb_lg	md	mh	m1	ms_20	mw	mbw	mwr	mwW
tsunami										
0	5.6	3.5	4.11	1.1	4.2	NaN	3.83	5.8	4.8	6.0
1	6.1	NaN	NaN	NaN	5.1	5.7	4.41	NaN	NaN	7.5

5. Calculate the rolling 60-day aggregations of OHLC data by ticker for the FAANG data. Use the same aggregations as exercise no. 3.

```
# Ensure the FAANG data is loaded and indexed by date
faang_df = pd.read_csv('/content/faang (1).csv', parse_dates=['date'])
faang_df.set_index('date', inplace=True)

# Calculate the rolling 60-day aggregations of OHLC data by ticker
faang_rolling_60 = faang_df.groupby('ticker').rolling(window='60D').agg({
    'open': 'mean',
    'high': 'max',
    'low': 'min',
    'close': 'mean',
    'volume': 'sum'
}).dropna()

faang_rolling_60.head()
```

		open	high	low	close	volume
ticker	date					
AAPL	2018-01-02	166.927100	169.0264	166.0442	168.987200	25555934.0
	2018-01-03	168.089600	171.2337	166.0442	168.972500	55073833.0
	2018-01-04	168.480367	171.2337	166.0442	169.229200	77508430.0
	2018-01-05	168.896475	172.0381	166.0442	169.840675	101168448.0
	2018-01-08	169.324680	172.2736	166.0442	170.080040	121736214.0

6. Create a pivot table of the FAANG data that compares the stocks. Put the ticker in the rows and show the averages of the OHLC and volume traded data.

```
# Create a pivot table for the FAANG data with ticker in the rows and averages of OHLC and volume
faang_pivot = faang_df.pivot_table(index='ticker', values=['open', 'high', 'low', 'close', 'volume'], aggfunc='mean')
```

		close	high	low	open	volume
ticker						
AAPL		186.986218	188.906858	185.135729	187.038674	3.402145e+07
AMZN		1641.726175	1662.839801	1619.840398	1644.072669	5.649563e+06
FB		171.510936	173.615298	169.303110	171.454424	2.768798e+07
GOOG		1113.225139	1125.777649	1101.001594	1113.554104	1.742645e+06
NFLX		319.290299	325.224583	313.187273	319.620533	1.147030e+07

7. Calculate the Z-scores for each numeric column of Netflix's data (ticker is NFLX) using apply().

```
# Filter the FAANG data for Netflix (NFLX) only
nflx_data = faang_df[faang_df['ticker'] == 'NFLX']

# Calculate Z-scores for each numeric column using apply()
nflx_z_scores = nflx_data.select_dtypes(include=['float64', 'int64']).apply(lambda x: (x - x.mean()) / x.std())

nflx_z_scores.head()
```

		open	high	low	close	volume
	date					
2018-01-02		-2.500753	-2.516023	-2.410226	-2.416644	-0.088760
2018-01-03		-2.380291	-2.423180	-2.285793	-2.335286	-0.507606
2018-01-04		-2.296272	-2.406077	-2.234616	-2.323429	-0.959287
2018-01-05		-2.275014	-2.345607	-2.202087	-2.234303	-0.782331
2018-01-08		-2.218934	-2.295113	-2.143759	-2.192192	-1.038531

8. Add event descriptions:

- Create a dataframe with the following three columns: ticker, date, and event. The columns should have the following values:
- ticker: 'FB'
- date: ['2018-07-25', '2018-03-19', '2018-03-20'] event: ['Disappointing user growth announced after close.', 'Cambridge Analytica story', 'FTC investigation']
- Set the index to ['date', 'ticker']
- Merge this data with the FAANG data using an outer join

```
# Create the events DataFrame
events_data = pd.DataFrame({
    'ticker': ['FB', 'FB', 'FB'],
    'date': ['2018-07-25', '2018-03-19', '2018-03-20'],
    'event': ['Disappointing user growth announced after close.', 'Cambridge Analytica story', 'FTC investigation']
})

# Convert 'date' to datetime format and set ['date', 'ticker'] as the index
events_data['date'] = pd.to_datetime(events_data['date'])
events_data.set_index(['date', 'ticker'], inplace=True)

# Ensure the FAANG data is in the correct format for merging
faang_df.reset_index(inplace=True)
faang_df.set_index(['date', 'ticker'], inplace=True)

# Merge the events data with the FAANG data using an outer join
faang_with_events = faang_df.join(events_data, on=['date', 'ticker'], how='outer').sort_index()

faang_with_events[faang_with_events['event'].notnull()]
```

		open	high	low	close	volume	event
	date ticker						
2018-03-19	FB	177.010	177.17	170.06	172.56	88140060	Cambridge Analytica story
2018-03-20	FB	167.470	170.20	161.95	168.15	129851768	FTC investigation
2018-07-25	FB	215.715	218.62	214.27	217.50	64592585	Disappointing user growth announced after close.

9. Use the transform() method on the FAANG data to represent all the values in terms of the first date in the data. To do so, divide all the values for each ticker by the values for the first date in the data for that ticker. This is referred to as an index, and the data for the first date

is the base ([https://ec.europa.eu/eurostat/statistics-explained/index.php/](https://ec.europa.eu/eurostat/statistics-explained/index.php/Beginners:_Statisticalconcept-Indexandbaseyear) Beginners: Statisticalconcept-Indexandbaseyear). When data is in this format, we can easily see growth over time. Hint: transform() can take a function name.

```
# Reset index of the FAANG data to work with transform()
faang_df.reset_index(inplace=True)
faang_df.set_index('date', inplace=True)

# Function to calculate the index relative to the first row in the group
def index_relative_to_first(row):
    return row / row.iloc[0]

# Group by ticker and apply the transformation to numeric columns
faang_indexed = faang_df.groupby('ticker').transform(index_relative_to_first)

# Include the 'ticker' column back into the faang_indexed DataFrame for clarity
faang_indexed['ticker'] = faang_df['ticker']

faang_indexed.head()
```

	open	high	low	close	volume	ticker
date						
2018-01-02	1.000000	1.000000	1.000000	1.000000	1.000000	FB
2018-01-03	1.023638	1.017623	1.021290	1.017914	0.930292	FB
2018-01-04	1.040635	1.025498	1.036889	1.016040	0.764707	FB
2018-01-05	1.044518	1.029298	1.041566	1.029931	0.747830	FB
2018-01-08	1.053579	1.040313	1.049451	1.037813	0.991341	FB

Conclusion

In conclusion for the FAANG (Facebook, Apple, Amazon, Netflix, Google) and earthquake datasets yielded significant findings, showcasing the power of data manipulation and analysis. In examining earthquake data, we identified specific quakes in Japan based on magnitude and type, revealing the frequency distribution of these events and their potential to cause tsunamis. This analysis highlighted the commonality of lower magnitude earthquakes and provided insights into the relationship between earthquake characteristics and tsunamis. For the FAANG stocks, our investigation into monthly trading behaviors uncovered trends in prices and volumes, enhanced by a rolling 60-day aggregation for a deeper view of medium-term trends. A pivot table comparing the FAANG stocks illustrated variations in trading behavior, while the application of Z-scores to Netflix's data offered a statistical comparison of its performance. Additionally, merging specific event data with the FAANG dataset linked stock performance to major events, enriching the analysis. Transforming FAANG data to index values based on the first date facilitated